

Model9

May 2, 2026

1 Model 9 - PPO + Qwen2.5-7B

1.0.1 Upgraded from 1.5B to 7B model + 12 new difficulty-detection features

1.0.2 Expected: higher accuracy ceiling + larger budget std (better per-question allocation)

1.0.3 Cell 1: Install (skip if done), restart kernel, then Cell 2

```
[1]: # Cell 1: Install
import subprocess, sys
subprocess.check_call([sys.executable, '-m', 'pip', 'install',
    'torch', '--index-url', 'https://download.pytorch.org/whl/cu128', '-q'])
subprocess.check_call([sys.executable, '-m', 'pip', 'install',
    'numpy', 'transformers', 'datasets', 'sentence-transformers', 'accelerate',
    '-q'])
import torch
print('torch:', torch.__version__)
print('cuda:', torch.cuda.is_available())
if torch.cuda.is_available(): print('gpu:', torch.cuda.get_device_name(0))
print('DONE - restart kernel then run Cell 2')
```

```
torch: 2.11.0+cu128
cuda: True
gpu: NVIDIA H200 NVL
DONE - restart kernel then run Cell 2
```

```
[2]: import os
import re
import json
import copy
import time
import random
import shutil
import hashlib
from dataclasses import dataclass
from typing import Dict, List, Tuple, Optional

import numpy as np
```

```

import torch
import torch.nn as nn
import torch.optim as optim

from datasets import load_dataset
from transformers import AutoTokenizer, AutoModelForCausalLM
from sentence_transformers import SentenceTransformer

# =====
# Config
# =====

@dataclass
class Config:
    # Dataset
    train_subset: int = 300
    test_subset: int = 100

    # Models
    llm_name: str = "Qwen/Qwen2.5-7B-Instruct"
    embed_name: str = "sentence-transformers/all-MiniLM-L6-v2"

    # PARALLELISM: how many questions to process at once on GPU
    # H200 150GB VRAM: 7B model ~15GB, batch=8 safe
    llm_batch_size: int = 8

    # Oracle budgets: only 128-220 where model actually works
    # Removed 96 and 112 because model fails 70-74% there
    oracle_budgets: Tuple[int, ...] = (
        128, 144, 160, 176, 192, 210, 220
    )

    # Supervised warm-start
    warmstart_epochs: int = 10
    warmstart_learning_rate: float = 5e-4

    # PPO specific
    rl_epochs: int = 10
    ppo_epochs: int = 2          # FIX: reduced from 4 to 2 - prevents
    ↪over-updating per rollout
    ppo_epsilon: float = 0.2     # clip range: ratio stays in [0.8, 1.2]
    ppo_batch_size: int = 32     # mini-batch size for PPO updates
    policy_learning_rate: float = 3e-4
    value_learning_rate: float = 8e-4
    entropy_coef: float = 0.02   # higher entropy for budget diversity with 7B
    value_loss_coef: float = 0.5

```

```

grad_clip_norm: float = 1.0

# Reward: rebalanced penalty so 176-token correct gets positive reward
# At 176 tokens:  $1.0 + 0.25 - 0.0005 \cdot 176 = 1.162$  (correct)
# At 128 tokens:  $1.0 + 0.25 - 0.0005 \cdot 128 = 1.186$  (correct)
# Small gap encourages efficiency without punishing longer budgets
reward_token_penalty: float = 0.0005 # 7B more accurate, moderate penalty
exact_match_bonus: float = 0.25

# Budget range: min raised to 128 (below this model fails reliably)
min_budget: int = 96
max_budget: int = 256

# Answer generation
answer_max_new_tokens: int = 20
do_sample: bool = False
temperature: float = 0.0

# RL warmup exploration in 160-220 where model succeeds
warmup_epochs_rl: int = 2
warmup_mix_rl: float = 0.25
warmup_min_budget_high: int = 160

# Oracle blend in PPO: supervised signal to keep deterministic head moving
use_oracle_blend: bool = True
oracle_blend_weight: float = 0.4 # increased: keeps deterministic head
↪ learning

# Hardware
seed: int = 42
device: str = "cuda" if torch.cuda.is_available() else "cpu"

# Logging
print_every: int = 20
debug_examples: int = 3
eval_print_every: int = 10

# Cache
cache_dir: str = "./cache_ppo_v15"
clear_cache_on_start: bool = True
oracle_cache_path: str = "oracle_labels_v15.json"

# Saved models
best_warmstart_path: str = "best_warmstart_v15.pt"
best_warmstart_acc_path: str = "best_warmstart_acc_v15.pt"
best_ppo_path: str = "best_ppo_v15.pt"
best_ppo_acc_path: str = "best_ppo_acc_v15.pt"

```

```

best_value_path: str = "best_value_v15.pt"
final_policy_path: str = "final_policy_v15.pt"
final_value_path: str = "final_value_v15.pt"

# Baselines
baseline_budgets: Tuple[int, ...] = (96, 128, 160, 192, 224, 256)

CFG = Config()

# =====
# Utilities
# =====

def set_seed(seed):
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    if torch.cuda.is_available():
        torch.cuda.manual_seed_all(seed)

def ensure_dir(path):
    os.makedirs(path, exist_ok=True)

def reset_cache_dir(cache_dir, retries=5, delay=0.5):
    if not os.path.exists(cache_dir):
        os.makedirs(cache_dir, exist_ok=True)
        return
    last_error = None
    for _ in range(retries):
        try:
            if os.path.exists(cache_dir):
                shutil.rmtree(cache_dir, ignore_errors=False)
            if os.path.exists(cache_dir):
                for root, dirs, files in os.walk(cache_dir, topdown=False):
                    for f in files:
                        try: os.remove(os.path.join(root, f))
                        except FileNotFoundError: pass
                    for d in dirs:
                        try: os.rmdir(os.path.join(root, d))
                        except OSError: pass
                if os.path.exists(cache_dir): os.rmdir(cache_dir)
            break
        except OSError as e:

```

```

        last_error = e
        time.sleep(delay)
    if os.path.exists(cache_dir):
        raise RuntimeError(f"Failed to clear cache: {cache_dir}") from _
↪last_error
    os.makedirs(cache_dir, exist_ok=True)

def clamp_budget(x, min_budget, max_budget):
    return int(max(min_budget, min(max_budget, round(x))))

def denormalize_budget(x, min_budget, max_budget):
    return x * (max_budget - min_budget) + min_budget

def budget_to_normalized(budget, min_budget, max_budget):
    return float(budget - min_budget) / float(max_budget - min_budget)

def clean_model_output(text):
    if text is None: return ""
    cleaned = text
    for marker in ["\nHuman:", "\nAssistant:", "\nUser:", "\nSystem:",
                  "Human:", "Assistant:", "User:", "System:",
                  "You are an AI assistant", "<|im_end|>", "<|endoftext|>"]:
        if marker in cleaned:
            cleaned = cleaned.split(marker)[0]
    return cleaned.strip()

# =====
# GSM8K helpers
# =====

def extract_gsm8k_final_answer(answer_text):
    if "####" in answer_text:
        return answer_text.split("####")[-1].strip()
    return answer_text.strip()

def normalize_numeric_string(text):
    if text is None: return ""
    text = text.strip().replace(",", "").replace("$", "").replace("%", "").
↪strip()
    matches = re.findall(r"-?\d+(?:\.\d+)?", text)
    if not matches: return text.lower().strip()

```

```

return matches[-1]

def parse_final_answer(output_text):
    text = clean_model_output(output_text)
    for pattern in [
        r"FINAL ANSWER\s*:\s*(-?\d+(?:\.\d+)?)",
        r"Final Answer\s*:\s*(-?\d+(?:\.\d+)?)",
        r"Answer\s*:\s*(-?\d+(?:\.\d+)?)",
    ]:
        m = re.search(pattern, text, re.IGNORECASE)
        if m: return normalize_numeric_string(m.group(1))
    lines = [l.strip() for l in text.splitlines() if l.strip()]
    first_line = lines[0] if lines else text.strip()
    m = re.search(r"-?\d+(?:\.\d+)?", first_line)
    if m: return normalize_numeric_string(m.group(0))
    matches = re.findall(r"-?\d+(?:\.\d+)?", text)
    if matches: return normalize_numeric_string(matches[0])
    return normalize_numeric_string(text)

def is_correct(pred, gold):
    return int(pred == gold)

def compute_reward(pred, gold, thinking_tokens, cfg):
    exact = float(is_correct(pred, gold))
    reward = exact - cfg.reward_token_penalty * float(thinking_tokens)
    if exact > 0.5:
        reward += cfg.exact_match_bonus
    return reward, exact

# =====
# Prompts
# =====

def build_thinking_prompt(question):
    return (
        "Solve the following grade-school math problem carefully.\n"
        "Reason step by step. Keep reasoning concise and relevant.\n"
        "Do not start a conversation. Only produce reasoning.\n\n"
        f"Question: {question}\n\nReasoning:"
    )

def build_answer_prompt(question, thinking_text):

```

```

    return (
        "Use the reasoning below to produce the final numeric answer.\n"
        "Return only: FINAL ANSWER: <number>\n\n"
        f"Question: {question}\n\nReasoning:\n{thinking_text}\n\nFINAL ANSWER:"
    )

# =====
# Caching
# =====

def make_cache_key(stage, question, extra, model_name):
    raw = f"{stage}|||{model_name}|||{question}|||{extra}"
    return hashlib.md5(raw.encode()).hexdigest()

def load_cache(cache_dir, key):
    path = os.path.join(cache_dir, f"{key}.json")
    if os.path.exists(path):
        with open(path, "r", encoding="utf-8") as f:
            return json.load(f)
    return None

def save_cache(cache_dir, key, value):
    path = os.path.join(cache_dir, f"{key}.json")
    with open(path, "w", encoding="utf-8") as f:
        json.dump(value, f, ensure_ascii=False, indent=2)

# =====
# LLM Reasoner with BATCHED GPU inference
# This is the main speedup vs Model 5
# Instead of 1 question at a time, processes batch_size questions
# simultaneously on the GPU - uses the H200 properly
# =====

class LLMReasoner:
    def __init__(self, model_name, device, batch_size=16):
        self.tokenizer = AutoTokenizer.from_pretrained(model_name)
        if self.tokenizer.pad_token is None:
            self.tokenizer.pad_token = self.tokenizer.eos_token

        # LEFT padding is required for decoder-only models in batched mode
        # Without this, the model attends to wrong positions
        self.tokenizer.padding_side = "left"

```

```

self.model = AutoModelForCausalLM.from_pretrained(
    model_name,
    dtype=torch.float16,
    device_map="auto"
)
self.model.eval()
self.device = device
self.batch_size = batch_size

if hasattr(self.model, "generation_config") and self.model.
↳ generation_config:
    self.model.generation_config.do_sample = False
    self.model.generation_config.temperature = None
    self.model.generation_config.top_p = None
    self.model.generation_config.top_k = None

def generate(self, prompt, max_new_tokens, do_sample=False, temperature=0.
↳ 0):
    """Single generation - keeps backward compatibility with cache system.
    ↳ """
    results = self.generate_batch([prompt], max_new_tokens, do_sample,
↳ temperature)
    return results[0]

def generate_batch(self, prompts, max_new_tokens, do_sample=False,
↳ temperature=0.0):
    """
    BATCHED GPU INFERENCE - the key parallelism speedup.

    How it works:
    - Takes a list of prompts
    - Tokenizes them all together with padding
    - GPU processes ALL of them in one forward pass
    - Returns results for all prompts

    Why this is faster:
    - Sequential: GPU runs at 5% utilization (mostly idle waiting)
    - Batched: GPU runs at 70-90% utilization
    - Speedup: 4-8x depending on batch size
    """
    all_results = []

    for i in range(0, len(prompts), self.batch_size):
        batch = prompts[i : i + self.batch_size]

        # Tokenize with padding so all sequences are same length
        inputs = self.tokenizer(

```



```

        batch,
        return_tensors="pt",
        padding=True,
        truncation=True,
        max_length=600,
    ).to(self.device)

    input_len = inputs["input_ids"].shape[1]

    gen_kwargs = {
        "max_new_tokens": max_new_tokens,
        "do_sample": do_sample,
        "pad_token_id": self.tokenizer.eos_token_id,
        "eos_token_id": self.tokenizer.eos_token_id,
        "attention_mask": inputs["attention_mask"],
    }
    if do_sample:
        gen_kwargs["temperature"] = temperature

    with torch.no_grad():
        # GPU processes entire batch at once
        outputs = self.model.generate(inputs["input_ids"], **gen_kwargs)

    for output in outputs:
        gen_ids = output[input_len:]
        text = clean_model_output(
            self.tokenizer.decode(gen_ids, skip_special_tokens=True)
        )
        all_results.append({"text": text, "tokens_used": len(gen_ids)})

    return all_results

# =====
# Two-stage example run (single, uses cache)
# =====

def run_two_stage_example(question, gold_answer, thinking_budget, llm, cfg):
    ensure_dir(cfg.cache_dir)

    # Stage 1: Thinking
    think_key = make_cache_key("think", question, f"budget={thinking_budget}",
    ↪cfg.llm_name)
    cached = load_cache(cfg.cache_dir, think_key)
    if cached:
        thinking_text = clean_model_output(cached["text"])
        thinking_tokens = cached["tokens_used"]

```

```

else:
    r = llm.generate(build_thinking_prompt(question), thinking_budget)
    thinking_text = clean_model_output(r["text"])
    thinking_tokens = r["tokens_used"]
    save_cache(cfg.cache_dir, think_key, {"text": thinking_text,
↪"tokens_used": thinking_tokens})

# Stage 2: Answer
answer_key = make_cache_key("answer", question, thinking_text, cfg.llm_name)
cached = load_cache(cfg.cache_dir, answer_key)
if cached:
    answer_text = clean_model_output(cached["text"])
    answer_tokens = cached["tokens_used"]
else:
    r = llm.generate(build_answer_prompt(question, thinking_text), cfg.
↪answer_max_new_tokens)
    answer_text = clean_model_output(r["text"])
    answer_tokens = r["tokens_used"]
    save_cache(cfg.cache_dir, answer_key, {"text": answer_text,
↪"tokens_used": answer_tokens})

pred = parse_final_answer(answer_text)
reward, exact = compute_reward(pred, gold_answer, thinking_tokens, cfg)
return {
    "prediction": pred, "gold": gold_answer, "correct": int(exact),
    "thinking_tokens_used": thinking_tokens, "answer_tokens_used":
↪answer_tokens,
    "reward": reward, "thinking_text": thinking_text, "answer_text":
↪answer_text,
    "thinking_budget": thinking_budget,
}

# =====
# BATCHED two-stage run for Oracle and RL rollouts
# Processes multiple questions at once using GPU parallelism
# =====

def run_batch_two_stage(questions, gold_answers, thinking_budget, llm, cfg):
    """
    Run two-stage inference for a BATCH of questions at the same budget.

    Step 1: Check cache for all questions
    Step 2: Run LLM on uncached questions in ONE batched GPU call
    Step 3: Run answer extraction in ONE batched GPU call
    Step 4: Save new results to cache

```

```

This replaces the sequential for-loop that caused the 4-hour runtime.
"""

ensure_dir(cfg.cache_dir)
n = len(questions)

# --- Stage 1: Thinking ---
thinking_texts = [None] * n
thinking_tokens_list = [None] * n
uncached_think_idx = []
uncached_think_prompts = []

# Check cache first for all questions
for i, q in enumerate(questions):
    key = make_cache_key("think", q, f"budget={thinking_budget}", cfg.
↳llm_name)
    cached = load_cache(cfg.cache_dir, key)
    if cached:
        thinking_texts[i] = clean_model_output(cached["text"])
        thinking_tokens_list[i] = cached["tokens_used"]
    else:
        uncached_think_idx.append(i)
        uncached_think_prompts.append(build_thinking_prompt(q))

# Batch GPU call for all uncached thinking
if uncached_think_prompts:
    batch_results = llm.generate_batch(uncached_think_prompts,
↳thinking_budget)
    for idx, result in zip(uncached_think_idx, batch_results):
        thinking_texts[idx] = clean_model_output(result["text"])
        thinking_tokens_list[idx] = result["tokens_used"]
        key = make_cache_key("think", questions[idx],
↳f"budget={thinking_budget}", cfg.llm_name)
        save_cache(cfg.cache_dir, key, {
            "text": thinking_texts[idx],
            "tokens_used": thinking_tokens_list[idx]
        })

# --- Stage 2: Answer extraction ---
answer_texts = [None] * n
answer_tokens_list = [None] * n
uncached_ans_idx = []
uncached_ans_prompts = []

for i in range(n):
    key = make_cache_key("answer", questions[i], thinking_texts[i], cfg.
↳llm_name)
    cached = load_cache(cfg.cache_dir, key)

```

```

        if cached:
            answer_texts[i] = clean_model_output(cached["text"])
            answer_tokens_list[i] = cached["tokens_used"]
        else:
            uncached_ans_idx.append(i)
            uncached_ans_prompts.append(
                build_answer_prompt(questions[i], thinking_texts[i])
            )

    # Batch GPU call for all uncached answers
    if uncached_ans_prompts:
        batch_results = llm.generate_batch(uncached_ans_prompts, cfg.
        ↪answer_max_new_tokens)
        for idx, result in zip(uncached_ans_idx, batch_results):
            answer_texts[idx] = clean_model_output(result["text"])
            answer_tokens_list[idx] = result["tokens_used"]
            key = make_cache_key("answer", questions[idx], thinking_texts[idx],
            ↪cfg.llm_name)
            save_cache(cfg.cache_dir, key, {
                "text": answer_texts[idx],
                "tokens_used": answer_tokens_list[idx]
            })

    # --- Compute rewards ---
    results = []
    for i in range(n):
        pred = parse_final_answer(answer_texts[i])
        reward, exact = compute_reward(pred, gold_answers[i],
        ↪thinking_tokens_list[i], cfg)
        results.append({
            "prediction": pred, "gold": gold_answers[i], "correct": int(exact),
            "thinking_tokens_used": thinking_tokens_list[i],
            "answer_tokens_used": answer_tokens_list[i],
            "reward": reward, "thinking_text": thinking_texts[i],
            "answer_text": answer_texts[i], "thinking_budget": thinking_budget,
        })
    return results

# =====
# Featurizer
# =====

class QuestionFeaturizer:
    """
    Rich difficulty-aware featurizer.

```

New: 12 difficulty features on top of original 50 handcrafted = 62 total,
↳ handcrafted.

Total input dim: 384 (SBERT all-MiniLM-L6-v2) + 62 = 446.

```

"""
def __init__(self, embed_model_name, device):
    self.embedder = SentenceTransformer(embed_model_name, device=device)

def handcrafted_features(self, question):
    q = question.lower()
    tokens = q.split()
    nums = re.findall(r"\d+\.?\\d*", q)

    # --- Original 50 features (preserved for backward compat) ---
    orig_keywords = [
        "total", "left", "remaining", "each", "every", "more", "less",
        "after", "before", "twice", "half", "times", "altogether",
        "difference", "sum", "product", "cost", "price", "sold",
        "per", "average", "equal", "share", "then", "gave", "bought",
        "earned", "minutes", "hours", "days", "weeks", "fraction",
        "ratio", "percent", "another", "still", "now", "together"
    ]
    orig_feats = [
        len(tokens),
        len(nums),
        len(q),
        max(1, len(re.findall(r"[.!?]", q))),
        1.0 if any(w in q for w in ["more", "sum", "altogether", "total",
↳ "together"]) else 0.0,
        1.0 if any(w in q for w in ["left", "remaining", "difference",
↳ "less", "still"]) else 0.0,
        1.0 if any(w in q for w in ["each", "every", "times", "product",
↳ "twice"]) else 0.0,
        1.0 if any(w in q for w in ["per", "average", "equally", "half",
↳ "share", "ratio"]) else 0.0,
        1.0 if any(w in q for w in ["then", "after", "before",
↳ "altogether", "remaining", "now"]) else 0.0,
        1.0 if any(w in q for w in ["minutes", "hours", "days", "weeks"])
↳ else 0.0,
        1.0 if "$" in q or any(w in q for w in ["dollar", "dollars",
↳ "cost", "price"]) else 0.0,
        1.0 if any(w in q for w in ["half", "fraction", "percent",
↳ "ratio"]) else 0.0,
    ] + [1.0 if kw in q else 0.0 for kw in orig_keywords]

    # --- NEW: 12 difficulty-detection features ---

```

```

# 1. Operation count - how many arithmetic operations are implied
op_words = [
    "add", "added", "plus", "more", "total", "sum", "altogether",
    ↪ "combined",
    "subtract", "subtracted", "minus", "less", "fewer", "remain",
    ↪ "left", "difference",
    "multiply", "multiplied", "times", "product", "double", "triple",
    ↪ "twice",
    "divide", "divided", "split", "share", "each", "per", "half",
    ↪ "quarter",
    "percent", "percentage", "fraction", "ratio", "rate"
]
op_count = float(sum(1 for w in tokens if w in op_words))

# 2. Numerical value count
num_count = float(len(nums))

# 3. Large number flag (>100 implies non-trivial arithmetic)
has_large_num = 1.0 if any(
    float(n) > 100 for n in nums
    if re.match(r"^\d+\.? \d*$", n)
) else 0.0

# 4. Sentence count - proxy for reasoning steps required
sentence_count = float(max(1, len(re.findall(r"[!?]", q)) + 1))

# 5. Multi-step sequence indicators
multistep_words = [
    "first", "then", "next", "after that", "finally", "lastly",
    "second", "third", "step", "now", "before", "once"
]
multistep_count = float(sum(1 for w in multistep_words if w in q))

# 6. Conditional logic (if/when makes questions harder)
has_conditional = 1.0 if any(w in q for w in [
    "if ", "when ", "unless", "until", "provided"
]) else 0.0

# 7. Comparison operators
has_comparison = 1.0 if any(w in q for w in [
    "more than", "less than", "greater than", "fewer than",
    "at least", "at most", "how many more", "how many fewer",
    "how much more", "how much less"
]) else 0.0

# 8. Unit conversion (time/distance/weight add steps)
has_unit_conversion = 1.0 if any(w in q for w in [

```

```

        "minute", "hour", "second", "day", "week", "month", "year",
        "inch", "foot", "feet", "yard", "mile", "meter", "kilometer",
        "ounce", "pound", "gram", "kilogram", "liter", "gallon"
    ]) else 0.0

    # 9. Percentage/fraction complexity
    has_pct_fraction = 1.0 if any(w in q for w in [
        "percent", "%", "fraction", "half", "third", "quarter",
        "one-half", "one-third", "two-thirds", "one-quarter"
    ]) else 0.0

    # 10. Entity count (people/objects to track)
    entity_indicators = [
        "he", "she", "they", "his", "her", "their",
        "john", "mary", "tom", "jane", "bob", "alice",
        "sam", "mike", "sarah", "david", "lisa", "james"
    ]
    entity_count = float(sum(1 for w in tokens if w in entity_indicators))

    # 11. Multiple sub-questions
    has_multiple_q = 1.0 if (
        q.count("?") > 1 or "and how" in q or "also how" in q
    ) else 0.0

    # 12. Composite difficulty score (normalized 0-1)
    raw_difficulty = (
        min(op_count, 8.0) / 8.0 * 0.35 +
        min(num_count, 6.0) / 6.0 * 0.25 +
        min(sentence_count, 5.0) / 5.0 * 0.20 +
        min(multistep_count, 4.0) / 4.0 * 0.20
    )

    new_feats = [
        op_count, num_count, has_large_num, sentence_count,
        multistep_count, has_conditional, has_comparison,
        has_unit_conversion, has_pct_fraction, entity_count,
        has_multiple_q, raw_difficulty,
    ]

    return np.array(orig_feats + new_feats, dtype=np.float32)

def encode(self, question):
    emb = self.embedder.encode(
        question, convert_to_numpy=True, normalize_embeddings=True
    ).astype(np.float32)
    return np.concatenate([emb, self.handcrafted_features(question)],
        ↪axis=0)

```

```

def encode_batch(self, questions):
    embs = self.embedder.encode(
        questions, convert_to_numpy=True,
        normalize_embeddings=True, batch_size=64
    ).astype(np.float32)
    hand = np.stack([self.handcrafted_features(q) for q in questions])
    return np.concatenate([embs, hand], axis=1)

# =====
# Policy + Value Networks
# =====

class ContinuousBudgetPolicy(nn.Module):
    def __init__(self, input_dim, hidden_dim=256):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(input_dim, hidden_dim), nn.ReLU(),
            nn.Linear(hidden_dim, hidden_dim), nn.ReLU(),
        )
        self.mean_head = nn.Linear(hidden_dim, 1)
        # log_std = 0.5 gives std ~1.65 for initial exploration
        self.log_std = nn.Parameter(torch.tensor([0.5], dtype=torch.float32))

    def forward(self, x):
        h = self.net(x)
        mean = self.mean_head(h)
        return mean, self.log_std.expand_as(mean)

    def sample_budget(self, x):
        mean, log_std = self.forward(x)
        dist = torch.distributions.Normal(mean, torch.exp(log_std))
        raw = dist.rsample()
        normalized = torch.sigmoid(raw)
        log_prob = dist.log_prob(raw).sum(dim=-1)
        entropy = dist.entropy().sum(dim=-1)
        return normalized, log_prob, entropy

    def get_log_prob(self, x, normalized_budgets):
        """
        Get log probability for given normalized budgets.
        Used in PPO to compute ratio between new and old policy.
        """
        mean, log_std = self.forward(x)
        dist = torch.distributions.Normal(mean, torch.exp(log_std))
        # Inverse sigmoid to get raw values

```



```

        raw = torch.logit(normalized_budgets.unsqueeze(-1).clamp(1e-6, 1 - 1e-6))
        log_prob = dist.log_prob(raw).sum(dim=-1)
        entropy = dist.entropy().sum(dim=-1)
        return log_prob, entropy

    def deterministic_budget(self, x):
        mean, _ = self.forward(x)
        return torch.sigmoid(mean)

class ValueNetwork(nn.Module):
    def __init__(self, input_dim, hidden_dim=256):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(input_dim, hidden_dim), nn.ReLU(),
            nn.Linear(hidden_dim, hidden_dim), nn.ReLU(),
            nn.Linear(hidden_dim, 1),
        )

    def forward(self, x):
        return self.net(x).squeeze(-1)

# =====
# Evaluation
# =====

def evaluate_fixed_budget(data, llm, cfg, fixed_budget):
    corrects, tokens, rewards = [], [], []
    questions = [item["question"] for item in data]
    gold_answers = [
        normalize_numeric_string(extract_gsm8k_final_answer(item["answer"]))
        for item in data
    ]
    print(f"\nEvaluating fixed budget = {fixed_budget}...")

    # Process in batches
    for i in range(0, len(questions), cfg.llm_batch_size):
        batch_q = questions[i : i + cfg.llm_batch_size]
        batch_g = gold_answers[i : i + cfg.llm_batch_size]
        results = run_batch_two_stage(batch_q, batch_g, fixed_budget, llm, cfg)
        for r in results:
            corrects.append(r["correct"])
            tokens.append(r["thinking_tokens_used"])
            rewards.append(r["reward"])
        done = min(i + cfg.llm_batch_size, len(questions))

```

```

        if done % cfg.eval_print_every == 0 or done == len(questions):
            print(f" Fixed {fixed_budget} | {done}/{len(questions)} | "
                  f"Acc={np.mean(corrects):.4f} | AvgReward={np.mean(rewards):.4f}")

    return {
        "accuracy": float(np.mean(corrects)),
        "avg_thinking_tokens": float(np.mean(tokens)),
        "avg_reward": float(np.mean(rewards)),
        "budget": fixed_budget
    }

def evaluate_policy_deterministic(data, featurizer, policy, llm, cfg):
    policy.eval()
    corrects, tokens, rewards, budgets = [], [], [], []

    questions = [item["question"] for item in data]
    gold_answers = [
        normalize_numeric_string(extract_gsm8k_final_answer(item["answer"]))
        for item in data
    ]

    # Encode all features at once
    all_feats = featurizer.encode_batch(questions)

    with torch.no_grad():
        feats_tensor = torch.tensor(all_feats, dtype=torch.float32, device=cfg.
        device)

        norm_budgets = policy.deterministic_budget(feats_tensor).squeeze(-1)
        budget_list = [
            clamp_budget(
                denormalize_budget(nb.item(), cfg.min_budget, cfg.max_budget),
                cfg.min_budget, cfg.max_budget
            )
            for nb in norm_budgets
        ]

    # Group questions by budget for batched inference
    from collections import defaultdict
    budget_groups = defaultdict(list)
    for i, b in enumerate(budget_list):
        budget_groups[b].append(i)

    results_by_idx = {}
    for budget, indices in budget_groups.items():
        batch_q = [questions[i] for i in indices]

```

```

        batch_g = [gold_answers[i] for i in indices]
        batch_results = run_batch_two_stage(batch_q, batch_g, budget, llm, cfg)
        for i, r in zip(indices, batch_results):
            results_by_idx[i] = r

    for i in range(len(data)):
        r = results_by_idx[i]
        corrects.append(r["correct"])
        tokens.append(r["thinking_tokens_used"])
        rewards.append(r["reward"])
        budgets.append(budget_list[i])

    if corrects:
        print(f" Det eval complete | Acc={np.mean(corrects):.4f} | "
              f"AvgBudget={np.mean(budgets):.1f} | StdBudget={np.std(budgets):.1f} | "
              f"Range=[{np.min(budgets):.0f}, {np.max(budgets):.0f}]")

    return {
        "accuracy": float(np.mean(corrects)),
        "avg_thinking_tokens": float(np.mean(tokens)),
        "avg_reward": float(np.mean(rewards)),
        "avg_budget": float(np.mean(budgets)),
        "min_budget": float(np.min(budgets)),
        "max_budget": float(np.max(budgets)),
        "std_budget": float(np.std(budgets)),
    }

# =====
# Oracle builder using BATCHED inference
# Model 5 took ~2.5 hours for this step alone
# With batching: processes 16 questions per budget in one GPU call
# Expected speedup: 3-5x (from ~2.5 hours to ~30-45 min)
# =====

def build_oracle_budget_labels(train_data, llm, cfg):
    if os.path.exists(cfg.oracle_cache_path):
        print(f"\nLoading oracle labels from {cfg.oracle_cache_path}...")
        with open(cfg.oracle_cache_path, "r", encoding="utf-8") as f:
            return json.load(f)

    print("\nBuilding oracle labels with BATCHED GPU inference...")
    questions = [item["question"] for item in train_data]
    gold_answers = [
        normalize_numeric_string(extract_gsm8k_final_answer(item["answer"]))
        for item in train_data
    ]

```

```

]
n = len(questions)

# Initialize records
oracle_records = [
    {"question": q, "gold_answer": g, "budget_rewards": {}, "any_correct": False}
    for q, g in zip(questions, gold_answers)
]

# For each budget, process ALL 300 questions in batches
# This is much faster than the nested loop in Model 5
for budget in cfg.oracle_budgets:
    print(f" Trying budget={budget} for all {n} questions...")
    t0 = time.time()

    for i in range(0, n, cfg.llm_batch_size):
        batch_q = questions[i : i + cfg.llm_batch_size]
        batch_g = gold_answers[i : i + cfg.llm_batch_size]
        batch_results = run_batch_two_stage(batch_q, batch_g, budget, llm, cfg)

        for j, result in enumerate(batch_results):
            idx = i + j
            oracle_records[idx]["budget_rewards"][str(budget)] = {
                "reward": float(result["reward"]),
                "correct": int(result["correct"]),
                "thinking_tokens_used": int(result["thinking_tokens_used"]),
                "prediction": result["prediction"],
            }
            if result["correct"] == 1:
                oracle_records[idx]["any_correct"] = True

    elapsed = time.time() - t0
    correct_at_budget = sum(
        1 for r in oracle_records
        if r["budget_rewards"].get(str(budget), {}).get("correct", 0) == 1
    )
    print(f" Budget {budget} done in {elapsed:.1f}s | "
          f"Correct: {correct_at_budget}/{n} ({100*correct_at_budget/n:.1f}%)")

# Assign best budget per question
# FIX from Model 5: use HIGHEST REWARD budget, not minimum correct
# Model 5 assigned 57% of questions to budget=96 which skewed everything
for record in oracle_records:
    br = record["budget_rewards"]

```

```

correct_budgets = [int(b) for b, v in br.items() if v["correct"] == 1]
if correct_budgets:
    best_budget = min(correct_budgets)
else:
    best_budget = int(max(br.keys(), key=lambda b: br[b]["reward"]))
record["best_budget"] = best_budget
record["best_reward"] = float(br[str(best_budget)]["reward"])

# Print distribution to verify it looks reasonable
best_budgets = [r["best_budget"] for r in oracle_records]
solvable = [r["any_correct"] for r in oracle_records]
print(f"\nOracle stats: mean={np.mean(best_budgets):.1f}, "
      f"std={np.std(best_budgets):.1f}, "
      f"solvable={np.mean(solvable):.3f}")
for b in sorted(set(best_budgets)):
    count = sum(1 for x in best_budgets if x == b)
    print(f"  budget={b}: {count} questions ({100*count/len(best_budgets):.1f}%)")

with open(cfg.oracle_cache_path, "w", encoding="utf-8") as f:
    json.dump(oracle_records, f, ensure_ascii=False, indent=2)
print(f"Saved oracle labels to {cfg.oracle_cache_path}")
return oracle_records

# =====
# Supervised warm-start
# =====

def warmstart_policy(oracle_records, test_data, featurizer, policy, llm, cfg):
    print("\nStarting supervised warm-start...")
    optimizer = optim.Adam(policy.parameters(), lr=cfg.warmstart_learning_rate)
    mse_loss = nn.MSELoss(reduction="none")

    all_budgets = np.array([r["best_budget"] for r in oracle_records], dtype=np.float32)
    budget_mean = float(all_budgets.mean())
    budget_std = float(all_budgets.std()) + 1e-8

    best_reward = -1e9
    best_accuracy = -1.0

    for epoch in range(cfg.warmstart_epochs):
        policy.train()
        random.shuffle(oracle_records)
        losses, pred_budgets, target_budgets = [], [], []

```

```

for step_idx, record in enumerate(oracle_records):
    target_budget = int(record["best_budget"])
    x = torch.tensor(
        featurizer.encode(record["question"]),
        dtype=torch.float32, device=cfg.device
    ).unsqueeze(0)

    pred_norm = policy.deterministic_budget(x)
    target_norm = torch.tensor(
        [[budget_to_normalized(target_budget, cfg.min_budget, cfg.
↪max_budget)]]],
        dtype=torch.float32, device=cfg.device
    )

    # Weight by reward quality + how much this question's budget
↪differs from mean
    reward_weight = max(0.1, float(record["best_reward"]))
    deviation_weight = 1.0 + abs(target_budget - budget_mean) /
↪budget_std
    loss = reward_weight * deviation_weight * mse_loss(pred_norm,
↪target_norm).mean()

    optimizer.zero_grad()
    loss.backward()
    torch.nn.utils.clip_grad_norm_(policy.parameters(), cfg.
↪grad_clip_norm)
    optimizer.step()

    losses.append(float(loss.item()))
    pred_budget = clamp_budget(
        denormalize_budget(pred_norm.item(), cfg.min_budget, cfg.
↪max_budget),
        cfg.min_budget, cfg.max_budget
    )
    pred_budgets.append(pred_budget)
    target_budgets.append(target_budget)

    if (step_idx + 1) % cfg.print_every == 0 or (step_idx + 1) ==
↪len(oracle_records):
        mae = np.mean(np.abs(np.array(pred_budgets) - np.
↪array(target_budgets)))
        print(f"[Warmstart {epoch+1}/{cfg.warmstart_epochs}] "
              f"Step {step_idx+1}/{len(oracle_records)} | "
              f"Loss={np.mean(losses):.4f} | "
              f"PredAvg={np.mean(pred_budgets):.1f} | ")

```

```

        f"TargetAvg={np.mean(target_budgets):.1f} | MAE={mae:.
↪2f}")

    print(f"Warmstart epoch {epoch+1} done | Loss={np.mean(losses):.4f}")
    eval_det = evaluate_policy_deterministic(test_data, featurizer, policy, llm,
↪cfg)

    if eval_det["avg_reward"] > best_reward:
        best_reward = eval_det["avg_reward"]
        torch.save(policy.state_dict(), cfg.best_warmstart_path)
        print(f" Saved best warmstart by reward = {best_reward:.4f}")

    if eval_det["accuracy"] > best_accuracy:
        best_accuracy = eval_det["accuracy"]
        torch.save(policy.state_dict(), cfg.best_warmstart_acc_path)
        print(f" Saved best warmstart by accuracy = {best_accuracy:.4f}")

# =====
# PPO Training
# This is the core fix over Model 5's basic policy gradient
#
# How PPO works here:
# 1. ROLLOUT: collect 300 (question, budget, reward) experiences
#           using current policy - all done with batched GPU calls
# 2. UPDATE: run 4 epochs of mini-batch updates using the PPO clip
#           The clip prevents any single update from being too large
# 3. REPEAT: for each RL epoch
#
# Why clip stops the oscillation:
# - Old code: loss = -(log_prob * advantage) + unbounded update size
# - PPO:      ratio = exp(new_log_prob - old_log_prob)
#           loss = -min(ratio*A, clip(ratio, 0.8, 1.2)*A)
#           If ratio > 1.2 or < 0.8, gradient is zeroed out
#           Policy cannot change by more than 20% per update
# =====

def train_ppo(train_data, test_data, featurizer, policy, value_net,
              llm, cfg, oracle_records=None):
    print("\nStarting PPO training...")
    print(f" PPO epsilon (clip range): {cfg.ppo_epsilon}")
    print(f" PPO epochs per rollout: {cfg.ppo_epochs}")
    print(f" Batch size: {cfg.llm_batch_size} questions at once on GPU")

    # Oracle lookup for blend signal
    oracle_lookup = {}
    if cfg.use_oracle_blend and oracle_records:

```

```

    for rec in oracle_records:
        oracle_lookup[rec["question"]] = rec["best_budget"]

    policy_optimizer = optim.Adam(policy.parameters(), lr=cfg.
    ↪policy_learning_rate)
    value_optimizer = optim.Adam(value_net.parameters(), lr=cfg.
    ↪value_learning_rate)
    mse_loss_fn = nn.MSELoss()

    questions = [item["question"] for item in train_data]
    gold_answers = [
        normalize_numeric_string(extract_gsm8k_final_answer(item["answer"]))
        for item in train_data
    ]

    # Pre-encode ALL training features once (saves time every epoch)
    print(" Pre-encoding all training features...")
    all_feats = featurizer.encode_batch(questions)
    all_feats_tensor = torch.tensor(all_feats, dtype=torch.float32, device=CFG.
    ↪device)

    best_test_reward = -1e9
    best_test_accuracy = -1.0

    for epoch in range(cfg.rl_epochs):
        print(f"\n{'='*80}")
        print(f"PP0 Epoch {epoch+1}/{cfg.rl_epochs}")
        print(f"{'='*80}")

        # -----
        # PHASE 1: ROLLOUT - collect experiences
        # Sample budgets for all questions, run LLM, get rewards
        # All LLM calls are batched for GPU parallelism
        # -----
        print(" Phase 1: Collecting rollout experiences (batched)...")
        policy.eval()

        rollout_idx = list(range(len(questions)))
        random.shuffle(rollout_idx)

        # Sample budgets for ALL questions at once
        with torch.no_grad():
            norm_budgets, old_log_probs, _ = policy.
            ↪sample_budget(all_feats_tensor)
            norm_budgets = norm_budgets.squeeze(-1)
            old_log_probs = old_log_probs.detach()

```



```

# Convert to actual budgets
budget_list = []
for i in range(len(questions)):
    nb = norm_budgets[i].item()
    b = clamp_budget(
        denormalize_budget(nb, cfg.min_budget, cfg.max_budget),
        cfg.min_budget, cfg.max_budget
    )
    # Warmup: mix in high budgets during first epochs
    if epoch < cfg.warmup_epochs_rl and random.random() < cfg.
↳warmup_mix_rl:
        b = random.randint(cfg.warmup_min_budget_high, cfg.max_budget)
        budget_list.append(b)

# Group questions by budget, run batched LLM inference
from collections import defaultdict
budget_groups = defaultdict(list)
for i, b in enumerate(budget_list):
    budget_groups[b].append(i)

rewards_list = [None] * len(questions)
corrects_list = [None] * len(questions)
tokens_list = [None] * len(questions)

for budget, indices in budget_groups.items():
    batch_q = [questions[i] for i in indices]
    batch_g = [gold_answers[i] for i in indices]
    batch_results = run_batch_two_stage(batch_q, batch_g, budget, llm,
↳cfg)

    for i, r in zip(indices, batch_results):
        rewards_list[i] = float(r["reward"])
        corrects_list[i] = int(r["correct"])
        tokens_list[i] = int(r["thinking_tokens_used"])

    rewards_tensor = torch.tensor(rewards_list, dtype=torch.float32,
↳device=cfg.device)

# Compute advantages
with torch.no_grad():
    values = value_net(all_feats_tensor)
    advantages = rewards_tensor - values
    # Normalize advantages for training stability
    advantages = (advantages - advantages.mean()) / (advantages.std() +
↳1e-8)

print(f" Rollout done | "
      f"Acc={np.mean(corrects_list):.4f} | ")

```

```

        f"AvgReward={np.mean(rewards_list):.4f} | "
        f"AvgBudget={np.mean(budget_list):.1f} | "
        f"StdBudget={np.std(budget_list):.1f}")

# -----
# PHASE 2: PPO UPDATE - update policy using clip
# Run ppo_epochs times on the same rollout data
# Mini-batches shuffled each time for better gradient estimates
# -----
print(f" Phase 2: PPO updates ({cfg.ppo_epochs} epochs, batch={cfg.
↪ppo_batch_size})...")
policy.train()

norm_budgets_stored = norm_budgets.detach()

for ppo_ep in range(cfg.ppo_epochs):
    indices = torch.randperm(len(questions))
    ep_policy_losses = []
    ep_value_losses = []
    ep_clip_fracs = []

    for start in range(0, len(questions), cfg.ppo_batch_size):
        batch_idx = indices[start : start + cfg.ppo_batch_size]

        batch_feats = all_feats_tensor[batch_idx]
        batch_old_log_probs = old_log_probs[batch_idx]
        batch_advantages = advantages[batch_idx]
        batch_rewards = rewards_tensor[batch_idx]
        batch_norm_budgets = norm_budgets_stored[batch_idx]

        # Get NEW log probs from updated policy
        new_log_probs, new_entropy = policy.get_log_prob(
            batch_feats, batch_norm_budgets
        )

        # PPO ratio: how much has the policy changed?
        ratio = torch.exp(new_log_probs - batch_old_log_probs)

        # THE PPO CLIP - this is the fix for the oscillation
        # ratio > 1.2: policy is going too far in a good direction, cap
↪it

        # ratio < 0.8: policy is going too far in a bad direction, cap
↪it

        clipped_ratio = torch.clamp(
            ratio,
            1.0 - cfg.ppo_epsilon,
            1.0 + cfg.ppo_epsilon

```

```

    )

    # Take the MINIMUM of clipped and unclipped
    # This is the pessimistic bound that prevents over-updating
    policy_loss = -torch.min(
        ratio * batch_advantages,
        clipped_ratio * batch_advantages
    ).mean()

    # Entropy bonus: keep policy from collapsing to one budget
    entropy_loss = -cfg.entropy_coef * new_entropy.mean()

    # Oracle blend: supervised signal on deterministic head
    # Keeps the deterministic head learning even during PPO
    supervised_loss = torch.tensor(0.0, device=cfg.device)
    if cfg.use_oracle_blend:
        oracle_targets = []
        valid_idx = []
        for k, bi in enumerate(batch_idx.tolist()):
            q = questions[bi]
            if q in oracle_lookup:
                oracle_targets.append(
                    budget_to_normalized(
                        oracle_lookup[q], cfg.min_budget, cfg.
↪max_budget
                    )
                )
                valid_idx.append(k)

        if oracle_targets:
            valid_feats = batch_feats[valid_idx]
            oracle_norm_t = torch.tensor(
                oracle_targets, dtype=torch.float32, device=cfg.
↪device
            ).unsqueeze(1)
            det_pred = policy.deterministic_budget(valid_feats)
            supervised_loss = nn.functional.mse_loss(det_pred, ↪
↪oracle_norm_t)

    total_policy_loss = (
        policy_loss + entropy_loss + cfg.oracle_blend_weight * ↪
↪supervised_loss
    )

    policy_optimizer.zero_grad()
    total_policy_loss.backward()

```

```

        torch.nn.utils.clip_grad_norm_(policy.parameters(), cfg.
↪grad_clip_norm)
        policy_optimizer.step()

        # Value network update
        value_preds = value_net(batch_feats)
        value_loss = mse_loss_fn(value_preds, batch_rewards)
        value_optimizer.zero_grad()
        (cfg.value_loss_coef * value_loss).backward()
        torch.nn.utils.clip_grad_norm_(value_net.parameters(), cfg.
↪grad_clip_norm)
        value_optimizer.step()

        # Track clip fraction (diagnostic: how often does clip activate?)
↪)
        clip_frac = ((ratio - 1.0).abs() > cfg.ppo_epsilon).float().
↪mean().item()

        ep_policy_losses.append(policy_loss.item())
        ep_value_losses.append(value_loss.item())
        ep_clip_fracs.append(clip_frac)

    print(f"    PPO update epoch {ppo_ep+1}/{cfg.ppo_epochs} | "
          f"PolicyLoss={np.mean(ep_policy_losses):.4f} | "
          f"ValueLoss={np.mean(ep_value_losses):.4f} | "
          f"ClipFrac={np.mean(ep_clip_fracs):.3f}")
    # ClipFrac guide:
    #   ~0.0: updates too small, increase LR or decrease epsilon
    #   0.1-0.3: healthy range
    #   >0.5: updates too large, something is wrong

    # Evaluate after each RL epoch
    print("\n Evaluating deterministic policy on test set...")
    eval_det = evaluate_policy_deterministic(test_data, featurizer, policy,
↪llm, cfg)
    print(f" Test Accuracy   : {eval_det['accuracy']:.4f}")
    print(f" Test AvgTokens   : {eval_det['avg_thinking_tokens']:.1f}")
    print(f" Test AvgReward    : {eval_det['avg_reward']:.4f}")
    print(f" Test Budget      : {eval_det['avg_budget']:.1f} +/- "
          f"{eval_det['std_budget']:.1f} "
          f"{eval_det['min_budget']:.0f}, {eval_det['max_budget']:.0f}")

    if eval_det["avg_reward"] > best_test_reward:
        best_test_reward = eval_det["avg_reward"]
        torch.save(policy.state_dict(), cfg.best_ppo_path)
        torch.save(value_net.state_dict(), cfg.best_value_path)
        print(f" Saved best PPO model by reward = {best_test_reward:.4f}")

```

```

        if eval_det["accuracy"] > best_test_accuracy:
            best_test_accuracy = eval_det["accuracy"]
            torch.save(policy.state_dict(), cfg.best_ppo_acc_path)
            print(f" Saved best PPO model by accuracy = {best_test_accuracy:.
↪4f}")

# =====
# Main
# =====

def main():
    set_seed(CFG.seed)

    print("torch version:", torch.__version__)
    print("cuda available:", torch.cuda.is_available())
    print("cuda device count:", torch.cuda.device_count())
    if torch.cuda.is_available():
        print("gpu name:", torch.cuda.get_device_name(0))
        print("gpu memory:", torch.cuda.get_device_properties(0).total_memory //
↪ 1e9, "GB")
    else:
        print("WARNING: No GPU detected. This will be very slow.")

    if CFG.clear_cache_on_start:
        print(f"\nClearing cache: {CFG.cache_dir}")
        reset_cache_dir(CFG.cache_dir)
        # Also remove any stale oracle cache from previous runs
        # This was the Model 6 bug: it loaded wrong oracle labels silently
        for stale in [CFG.oracle_cache_path,
↪ "oracle_labels_v11.json", "oracle_labels_v12.json",
↪ "oracle_labels_v13.json"]:
            if os.path.exists(stale):
                os.remove(stale)
                print(f" Removed stale oracle cache: {stale}")
            else:
                ensure_dir(CFG.cache_dir)

    print(f"Using device: {CFG.device}")
    print(f"LLM batch size: {CFG.llm_batch_size} (questions processed in
↪ parallel)")

    print("\nLoading GSM8K...")
    ds = load_dataset("gsm8k", "main")
    train_data = list(ds["train"].select(range(min(CFG.train_subset,
↪ len(ds["train"]))))))

```

```

    test_data = list(ds["test"].select(range(min(CFG.test_subset,
↪len(ds["test"]))))))
    print(f"Train: {len(train_data)} | Test: {len(test_data)}")

    print("\nLoading featurizer...")
    featurizer = QuestionFeaturizer(CFG.embed_name, CFG.device)
    input_dim = featurizer.encode(train_data[0]["question"]).shape[0]
    print(f"Feature dim: {input_dim}")

    print("\nLoading policy and value network...")
    policy = ContinuousBudgetPolicy(input_dim=input_dim).to(CFG.device)
    value_net = ValueNetwork(input_dim=input_dim).to(CFG.device)

    # Initialize toward 176 tokens (normalized ~0.72 for range 128-220)
    with torch.no_grad():
        policy.mean_head.bias.fill_(0.85)

    print("\nLoading LLM (with batched inference)...")
    llm = LLMReasoner(CFG.llm_name, CFG.device, batch_size=CFG.llm_batch_size)

    # Fixed budget baselines
    print("\n" + "="*70)
    print("FIXED BUDGET BASELINES")
    print("="*70)
    baseline_results = {}
    for b in CFG.baseline_budgets:
        res = evaluate_fixed_budget(test_data, llm, CFG, b)
        baseline_results[b] = res
        print(f"Fixed {b:>3} | Acc={res['accuracy']:.4f} | "
              f"AvgTokens={res['avg_thinking_tokens']:.0f} | "
              f"AvgReward={res['avg_reward']:.4f}")

    # Build oracle labels
    oracle_records = build_oracle_budget_labels(train_data, llm, CFG)

    # Warm-start
    warmstart_policy(oracle_records, test_data, featurizer, policy, llm, CFG)

    # Load best warm-start
    if os.path.exists(CFG.best_warmstart_acc_path):
        print("\nLoading best warm-start policy by accuracy...")
        policy.load_state_dict(torch.load(CFG.best_warmstart_acc_path,
↪map_location=CFG.device))
    elif os.path.exists(CFG.best_warmstart_path):
        print("\nLoading best warm-start policy by reward...")
        policy.load_state_dict(torch.load(CFG.best_warmstart_path,
↪map_location=CFG.device))

```

```

print("\nPost warm-start evaluation:")
warm_eval = evaluate_policy_deterministic(test_data, featurizer, policy,
llm, CFG)
print(json.dumps({"warmstart_eval": warm_eval}, indent=2))

# Pre-compute all training features (used for value pre-training and PPO)
print("\nPre-computing all training features...")
all_feats = featurizer.encode_batch([item["question"] for item in
train_data])
all_feats_tensor = torch.tensor(all_feats, dtype=torch.float32, device=CFG.
device)

# -----
# FIX: Pre-train value network before PPO starts
# Without this, ValueLoss starts at 2.6 (random init)
# causing garbage advantage estimates in PPO epoch 1
# Pre-training brings it down so epoch 1 updates are clean
# -----
print("\nPre-training value network on oracle rewards...")
value_pretrain_optimizer = optim.Adam(value_net.parameters(), lr=CFG.
value_learning_rate)
oracle_avg_reward = float(np.mean([r["best_reward"] for r in
oracle_records]))
print(f" Target oracle avg reward: {oracle_avg_reward:.4f}")

value_net.train()
for pretrain_ep in range(5):
    # Shuffle for better gradient estimates
    idx_perm = torch.randperm(len(train_data))
    ep_losses = []
    for start in range(0, len(train_data), 32):
        batch_idx = idx_perm[start:start+32]
        batch_feats = all_feats_tensor[batch_idx]

        # Per-question target: use oracle reward if available, else avg
        batch_targets = []
        for bi in batch_idx.tolist():
            q = train_data[bi]["question"]
            rec = next((r for r in oracle_records if r["question"] == q),
None)

            target = float(rec["best_reward"]) if rec else oracle_avg_reward
            batch_targets.append(target)

        target_t = torch.tensor(batch_targets, dtype=torch.float32,
device=CFG.device)

```

```

        preds = value_net(batch_feats)
        loss = nn.functional.mse_loss(preds, target_t)
        value_pretrain_optimizer.zero_grad()
        loss.backward()
        torch.nn.utils.clip_grad_norm_(value_net.parameters(), CFG.
↪grad_clip_norm)
        value_pretrain_optimizer.step()
        ep_losses.append(loss.item())

    print(f" Value pretrain epoch {pretrain_ep+1}/5 | Loss={np.
↪mean(ep_losses):.4f}")

    print("Value pre-training complete. PPO will start with better advantage_
↪estimates.\n")

# PPO Training
train_ppo(
    train_data=train_data,
    test_data=test_data,
    featurizer=featurizer,
    policy=policy,
    value_net=value_net,
    llm=llm,
    cfg=CFG,
    oracle_records=oracle_records,
)

# Load best PPO model
print("\nLoading best PPO model for final evaluation...")
if os.path.exists(CFG.best_ppo_acc_path):
    policy.load_state_dict(torch.load(CFG.best_ppo_acc_path,
↪map_location=CFG.device))
elif os.path.exists(CFG.best_ppo_path):
    policy.load_state_dict(torch.load(CFG.best_ppo_path, map_location=CFG.
↪device))
if os.path.exists(CFG.best_value_path):
    value_net.load_state_dict(torch.load(CFG.best_value_path,
↪map_location=CFG.device))

print("\nFinal deterministic policy evaluation:")
final_det = evaluate_policy_deterministic(test_data, featurizer, policy,
↪llm, CFG)
print(json.dumps({"final_deterministic_eval": final_det}, indent=2))

torch.save(policy.state_dict(), CFG.final_policy_path)
torch.save(value_net.state_dict(), CFG.final_value_path)

```



```

print(f"\nSaved: {CFG.final_policy_path}")
print(f"Saved: {CFG.final_value_path}")

# Final comparison table
print("\n" + "="*72)
print("FINAL RESULTS COMPARISON")
print("="*72)
print(f"{'Method':<35} {'Acc':>6} {'AvgTokens':>10} {'AvgReward':>10}")
print("-"*72)
for b, res in baseline_results.items():
    print(f"  Fixed budget {b:<22} {res['accuracy']:>6.4f} "
          f"{res['avg_thinking_tokens']:>10.0f} {res['avg_reward']:>10.4f}")
print(f"  Policy (PPO deterministic)      {final_det['accuracy']:>6.4f} "
      f"{final_det['avg_thinking_tokens']:>10.0f} {final_det['avg_reward']:"
      f">10.4f}")
print("="*72)
print("\nAI Disclosure: Written with Claude's assistance")

main()

```

/opt/conda/lib/python3.13/site-packages/tqdm/auto.py:21: TqdmWarning: IPProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html

```
from .autonotebook import tqdm as notebook_tqdm
```

torch version: 2.11.0+cu128

cuda available: True

cuda device count: 1

gpu name: NVIDIA H200 NVL

gpu memory: 150.0 GB

Clearing cache: ./cache_ppo_v15

Using device: cuda

LLM batch size: 8 (questions processed in parallel)

Loading GSM8K...

Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.

Train: 300 | Test: 100

Loading featurizer...

Loading weights: 100% | 103/103 [00:00<00:00, 12236.95it/s]

BertModel LOAD REPORT from: sentence-transformers/all-MiniLM-L6-v2

Key	Status	
embeddings.position_ids	UNEXPECTED	

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture;
not ok if you expect identical arch.

Feature dim: 446

Loading policy and value network...

Loading LLM (with batched inference)...

Fetching 4 files: 100%| | 4/4 [01:22<00:00, 20.74s/it]

Loading weights: 100%| | 339/339 [00:39<00:00, 8.52it/s]

=====

FIXED BUDGET BASELINES

=====

Evaluating fixed budget = 96...

Fixed 96 | 40/100 | Acc=0.4750 | AvgReward=0.5457

Fixed 96 | 80/100 | Acc=0.4625 | AvgReward=0.5301

Fixed 96 | 100/100 | Acc=0.4700 | AvgReward=0.5395

Fixed 96 | Acc=0.4700 | AvgTokens=96 | AvgReward=0.5395

Evaluating fixed budget = 128...

Fixed 128 | 40/100 | Acc=0.5500 | AvgReward=0.6235

Fixed 128 | 80/100 | Acc=0.6250 | AvgReward=0.7172

Fixed 128 | 100/100 | Acc=0.6300 | AvgReward=0.7235

Fixed 128 | Acc=0.6300 | AvgTokens=128 | AvgReward=0.7235

Evaluating fixed budget = 160...

Fixed 160 | 40/100 | Acc=0.6750 | AvgReward=0.7637

Fixed 160 | 80/100 | Acc=0.7125 | AvgReward=0.8106

Fixed 160 | 100/100 | Acc=0.7400 | AvgReward=0.8450

Fixed 160 | Acc=0.7400 | AvgTokens=160 | AvgReward=0.8450

Evaluating fixed budget = 192...

Fixed 192 | 40/100 | Acc=0.6750 | AvgReward=0.7477

Fixed 192 | 80/100 | Acc=0.7375 | AvgReward=0.8259

Fixed 192 | 100/100 | Acc=0.7600 | AvgReward=0.8540

Fixed 192 | Acc=0.7600 | AvgTokens=192 | AvgReward=0.8540

Evaluating fixed budget = 224...

Fixed 224 | 40/100 | Acc=0.7500 | AvgReward=0.8255

Fixed 224 | 80/100 | Acc=0.7875 | AvgReward=0.8724

Fixed 224 | 100/100 | Acc=0.7900 | AvgReward=0.8755

Fixed 224 | Acc=0.7900 | AvgTokens=224 | AvgReward=0.8755

Evaluating fixed budget = 256...
Fixed 256 | 40/100 | Acc=0.8250 | AvgReward=0.9032
Fixed 256 | 80/100 | Acc=0.8625 | AvgReward=0.9501
Fixed 256 | 100/100 | Acc=0.8400 | AvgReward=0.9220
Fixed 256 | Acc=0.8400 | AvgTokens=256 | AvgReward=0.9220

Building oracle labels with BATCHED GPU inference...
Trying budget=128 for all 300 questions...
Budget 128 done in 103.5s | Correct: 199/300 (66.3%)
Trying budget=144 for all 300 questions...
Budget 144 done in 90.7s | Correct: 214/300 (71.3%)
Trying budget=160 for all 300 questions...
Budget 160 done in 100.0s | Correct: 220/300 (73.3%)
Trying budget=176 for all 300 questions...
Budget 176 done in 103.9s | Correct: 238/300 (79.3%)
Trying budget=192 for all 300 questions...
Budget 192 done in 113.7s | Correct: 247/300 (82.3%)
Trying budget=210 for all 300 questions...
Budget 210 done in 127.3s | Correct: 255/300 (85.0%)
Trying budget=220 for all 300 questions...
Budget 220 done in 132.3s | Correct: 263/300 (87.7%)

Oracle stats: mean=140.2, std=23.9, solvable=0.947
budget=128: 215 questions (71.7%)
budget=144: 32 questions (10.7%)
budget=160: 10 questions (3.3%)
budget=176: 17 questions (5.7%)
budget=192: 13 questions (4.3%)
budget=210: 3 questions (1.0%)
budget=220: 10 questions (3.3%)

Saved oracle labels to oracle_labels_v15.json

Starting supervised warm-start...
[Warmstart 1/10] Step 20/300 | Loss=0.1758 | PredAvg=139.3 | TargetAvg=138.4 | MAE=37.65
[Warmstart 1/10] Step 40/300 | Loss=0.1080 | PredAvg=136.8 | TargetAvg=136.0 | MAE=27.98
[Warmstart 1/10] Step 60/300 | Loss=0.0831 | PredAvg=135.0 | TargetAvg=134.1 | MAE=25.23
[Warmstart 1/10] Step 80/300 | Loss=0.1079 | PredAvg=132.0 | TargetAvg=136.4 | MAE=27.59
[Warmstart 1/10] Step 100/300 | Loss=0.1157 | PredAvg=134.8 | TargetAvg=137.4 | MAE=27.91
[Warmstart 1/10] Step 120/300 | Loss=0.1089 | PredAvg=134.4 | TargetAvg=137.1 | MAE=27.27
[Warmstart 1/10] Step 140/300 | Loss=0.1479 | PredAvg=134.5 | TargetAvg=138.5 | MAE=28.80
[Warmstart 1/10] Step 160/300 | Loss=0.1557 | PredAvg=132.9 | TargetAvg=139.0 |

MAE=29.04
 [Warmstart 1/10] Step 180/300 | Loss=0.1539 | PredAvg=135.2 | TargetAvg=139.0 |
 MAE=29.86
 [Warmstart 1/10] Step 200/300 | Loss=0.1447 | PredAvg=135.1 | TargetAvg=139.0 |
 MAE=29.00
 [Warmstart 1/10] Step 220/300 | Loss=0.1634 | PredAvg=135.7 | TargetAvg=140.0 |
 MAE=30.01
 [Warmstart 1/10] Step 240/300 | Loss=0.1585 | PredAvg=135.8 | TargetAvg=140.2 |
 MAE=29.96
 [Warmstart 1/10] Step 260/300 | Loss=0.1555 | PredAvg=136.0 | TargetAvg=140.2 |
 MAE=29.76
 [Warmstart 1/10] Step 280/300 | Loss=0.1595 | PredAvg=135.3 | TargetAvg=140.2 |
 MAE=29.50
 [Warmstart 1/10] Step 300/300 | Loss=0.1616 | PredAvg=135.2 | TargetAvg=140.2 |
 MAE=29.23
 Warmstart epoch 1 done | Loss=0.1616
 Det eval complete | Acc=0.6500 | AvgBudget=136.1 | StdBudget=17.7 |
 Range=[101, 173]
 Saved best warmstart by reward = 0.7444
 Saved best warmstart by accuracy = 0.6500
 [Warmstart 2/10] Step 20/300 | Loss=0.2351 | PredAvg=140.1 | TargetAvg=142.2 |
 MAE=31.25
 [Warmstart 2/10] Step 40/300 | Loss=0.2094 | PredAvg=137.0 | TargetAvg=143.9 |
 MAE=30.07
 [Warmstart 2/10] Step 60/300 | Loss=0.1620 | PredAvg=137.3 | TargetAvg=138.9 |
 MAE=29.43
 [Warmstart 2/10] Step 80/300 | Loss=0.1309 | PredAvg=134.2 | TargetAvg=137.2 |
 MAE=26.73
 [Warmstart 2/10] Step 100/300 | Loss=0.1382 | PredAvg=134.8 | TargetAvg=138.0 |
 MAE=26.26
 [Warmstart 2/10] Step 120/300 | Loss=0.1356 | PredAvg=135.2 | TargetAvg=138.7 |
 MAE=25.68
 [Warmstart 2/10] Step 140/300 | Loss=0.1217 | PredAvg=135.6 | TargetAvg=137.8 |
 MAE=24.82
 [Warmstart 2/10] Step 160/300 | Loss=0.1203 | PredAvg=134.3 | TargetAvg=137.3 |
 MAE=24.44
 [Warmstart 2/10] Step 180/300 | Loss=0.1170 | PredAvg=135.2 | TargetAvg=137.9 |
 MAE=24.47
 [Warmstart 2/10] Step 200/300 | Loss=0.1314 | PredAvg=135.9 | TargetAvg=139.4 |
 MAE=25.45
 [Warmstart 2/10] Step 220/300 | Loss=0.1278 | PredAvg=135.3 | TargetAvg=139.2 |
 MAE=25.07
 [Warmstart 2/10] Step 240/300 | Loss=0.1351 | PredAvg=135.7 | TargetAvg=140.2 |
 MAE=24.95
 [Warmstart 2/10] Step 260/300 | Loss=0.1333 | PredAvg=137.1 | TargetAvg=140.6 |
 MAE=25.16
 [Warmstart 2/10] Step 280/300 | Loss=0.1339 | PredAvg=136.6 | TargetAvg=140.6 |
 MAE=25.16

[Warmstart 2/10] Step 300/300 | Loss=0.1273 | PredAvg=136.3 | TargetAvg=140.2 | MAE=24.37
 Warmstart epoch 2 done | Loss=0.1273
 Det eval complete | Acc=0.6400 | AvgBudget=135.0 | StdBudget=15.5 |
 Range=[102, 167]
 [Warmstart 3/10] Step 20/300 | Loss=0.0500 | PredAvg=141.3 | TargetAvg=137.6 | MAE=18.35
 [Warmstart 3/10] Step 40/300 | Loss=0.0590 | PredAvg=135.5 | TargetAvg=136.8 | MAE=18.32
 [Warmstart 3/10] Step 60/300 | Loss=0.0896 | PredAvg=134.9 | TargetAvg=137.8 | MAE=19.52
 [Warmstart 3/10] Step 80/300 | Loss=0.1390 | PredAvg=134.2 | TargetAvg=140.3 | MAE=23.36
 [Warmstart 3/10] Step 100/300 | Loss=0.1268 | PredAvg=136.0 | TargetAvg=139.9 | MAE=23.07
 [Warmstart 3/10] Step 120/300 | Loss=0.1149 | PredAvg=135.2 | TargetAvg=139.7 | MAE=22.43
 [Warmstart 3/10] Step 140/300 | Loss=0.1206 | PredAvg=136.3 | TargetAvg=139.8 | MAE=22.47
 [Warmstart 3/10] Step 160/300 | Loss=0.1109 | PredAvg=135.4 | TargetAvg=138.6 | MAE=21.64
 [Warmstart 3/10] Step 180/300 | Loss=0.1146 | PredAvg=135.0 | TargetAvg=139.1 | MAE=21.89
 [Warmstart 3/10] Step 200/300 | Loss=0.1165 | PredAvg=136.3 | TargetAvg=139.8 | MAE=22.36
 [Warmstart 3/10] Step 220/300 | Loss=0.1176 | PredAvg=135.6 | TargetAvg=139.6 | MAE=22.05
 [Warmstart 3/10] Step 240/300 | Loss=0.1237 | PredAvg=134.9 | TargetAvg=139.4 | MAE=21.65
 [Warmstart 3/10] Step 260/300 | Loss=0.1199 | PredAvg=135.2 | TargetAvg=140.1 | MAE=21.65
 [Warmstart 3/10] Step 280/300 | Loss=0.1194 | PredAvg=136.2 | TargetAvg=140.0 | MAE=22.17
 [Warmstart 3/10] Step 300/300 | Loss=0.1205 | PredAvg=135.4 | TargetAvg=140.2 | MAE=22.23
 Warmstart epoch 3 done | Loss=0.1205
 Det eval complete | Acc=0.6900 | AvgBudget=145.9 | StdBudget=10.7 |
 Range=[117, 171]
 Saved best warmstart by reward = 0.7896
 Saved best warmstart by accuracy = 0.6900
 [Warmstart 4/10] Step 20/300 | Loss=0.1631 | PredAvg=132.8 | TargetAvg=139.8 | MAE=25.45
 [Warmstart 4/10] Step 40/300 | Loss=0.1541 | PredAvg=133.1 | TargetAvg=140.6 | MAE=23.07
 [Warmstart 4/10] Step 60/300 | Loss=0.1489 | PredAvg=135.0 | TargetAvg=142.3 | MAE=24.12
 [Warmstart 4/10] Step 80/300 | Loss=0.1526 | PredAvg=134.5 | TargetAvg=144.2 | MAE=24.31

[Warmstart 4/10] Step 100/300 | Loss=0.1491 | PredAvg=134.5 | TargetAvg=143.8 |
 MAE=23.05
 [Warmstart 4/10] Step 120/300 | Loss=0.1386 | PredAvg=133.9 | TargetAvg=142.6 |
 MAE=21.46
 [Warmstart 4/10] Step 140/300 | Loss=0.1196 | PredAvg=133.5 | TargetAvg=140.8 |
 MAE=19.68
 [Warmstart 4/10] Step 160/300 | Loss=0.1249 | PredAvg=133.4 | TargetAvg=141.0 |
 MAE=19.69
 [Warmstart 4/10] Step 180/300 | Loss=0.1203 | PredAvg=132.7 | TargetAvg=140.7 |
 MAE=19.38
 [Warmstart 4/10] Step 200/300 | Loss=0.1167 | PredAvg=134.1 | TargetAvg=140.7 |
 MAE=19.88
 [Warmstart 4/10] Step 220/300 | Loss=0.1254 | PredAvg=133.5 | TargetAvg=140.7 |
 MAE=20.09
 [Warmstart 4/10] Step 240/300 | Loss=0.1238 | PredAvg=133.1 | TargetAvg=140.4 |
 MAE=19.61
 [Warmstart 4/10] Step 260/300 | Loss=0.1268 | PredAvg=133.2 | TargetAvg=141.0 |
 MAE=20.12
 [Warmstart 4/10] Step 280/300 | Loss=0.1181 | PredAvg=132.8 | TargetAvg=140.1 |
 MAE=19.20
 [Warmstart 4/10] Step 300/300 | Loss=0.1215 | PredAvg=132.9 | TargetAvg=140.2 |
 MAE=19.16
 Warmstart epoch 4 done | Loss=0.1215
 Det eval complete | Acc=0.6200 | AvgBudget=128.5 | StdBudget=7.6 | Range=[110,
 153]
 [Warmstart 5/10] Step 20/300 | Loss=0.0038 | PredAvg=127.0 | TargetAvg=128.8 |
 MAE=5.35
 [Warmstart 5/10] Step 40/300 | Loss=0.0470 | PredAvg=130.5 | TargetAvg=132.8 |
 MAE=11.12
 [Warmstart 5/10] Step 60/300 | Loss=0.0782 | PredAvg=132.1 | TargetAvg=136.0 |
 MAE=14.65
 [Warmstart 5/10] Step 80/300 | Loss=0.1029 | PredAvg=131.1 | TargetAvg=137.7 |
 MAE=15.85
 [Warmstart 5/10] Step 100/300 | Loss=0.0913 | PredAvg=131.9 | TargetAvg=137.7 |
 MAE=16.23
 [Warmstart 5/10] Step 120/300 | Loss=0.0920 | PredAvg=133.6 | TargetAvg=138.7 |
 MAE=17.00
 [Warmstart 5/10] Step 140/300 | Loss=0.0916 | PredAvg=133.7 | TargetAvg=138.9 |
 MAE=17.29
 [Warmstart 5/10] Step 160/300 | Loss=0.1002 | PredAvg=133.9 | TargetAvg=140.4 |
 MAE=17.92
 [Warmstart 5/10] Step 180/300 | Loss=0.0926 | PredAvg=133.8 | TargetAvg=139.8 |
 MAE=17.38
 [Warmstart 5/10] Step 200/300 | Loss=0.1012 | PredAvg=134.6 | TargetAvg=140.8 |
 MAE=18.68
 [Warmstart 5/10] Step 220/300 | Loss=0.0987 | PredAvg=135.2 | TargetAvg=140.8 |
 MAE=18.45
 [Warmstart 5/10] Step 240/300 | Loss=0.1093 | PredAvg=134.8 | TargetAvg=141.0 |

MAE=18.84
 [Warmstart 5/10] Step 260/300 | Loss=0.1051 | PredAvg=134.4 | TargetAvg=140.7 |
 MAE=18.35
 [Warmstart 5/10] Step 280/300 | Loss=0.1047 | PredAvg=133.9 | TargetAvg=140.4 |
 MAE=18.07
 [Warmstart 5/10] Step 300/300 | Loss=0.1030 | PredAvg=134.1 | TargetAvg=140.2 |
 MAE=18.11
 Warmstart epoch 5 done | Loss=0.1030
 Det eval complete | Acc=0.5600 | AvgBudget=119.2 | StdBudget=6.5 | Range=[106,
 142]
 [Warmstart 6/10] Step 20/300 | Loss=0.1743 | PredAvg=128.3 | TargetAvg=140.7 |
 MAE=16.55
 [Warmstart 6/10] Step 40/300 | Loss=0.1072 | PredAvg=129.4 | TargetAvg=137.6 |
 MAE=13.68
 [Warmstart 6/10] Step 60/300 | Loss=0.1396 | PredAvg=129.9 | TargetAvg=140.5 |
 MAE=16.42
 [Warmstart 6/10] Step 80/300 | Loss=0.1074 | PredAvg=130.7 | TargetAvg=138.6 |
 MAE=13.93
 [Warmstart 6/10] Step 100/300 | Loss=0.1190 | PredAvg=132.1 | TargetAvg=140.5 |
 MAE=15.38
 [Warmstart 6/10] Step 120/300 | Loss=0.1171 | PredAvg=133.9 | TargetAvg=140.5 |
 MAE=16.37
 [Warmstart 6/10] Step 140/300 | Loss=0.1100 | PredAvg=133.7 | TargetAvg=140.6 |
 MAE=16.26
 [Warmstart 6/10] Step 160/300 | Loss=0.1007 | PredAvg=133.5 | TargetAvg=140.2 |
 MAE=15.78
 [Warmstart 6/10] Step 180/300 | Loss=0.1104 | PredAvg=133.1 | TargetAvg=140.6 |
 MAE=16.18
 [Warmstart 6/10] Step 200/300 | Loss=0.1014 | PredAvg=132.7 | TargetAvg=139.6 |
 MAE=15.41
 [Warmstart 6/10] Step 220/300 | Loss=0.0994 | PredAvg=132.7 | TargetAvg=139.7 |
 MAE=15.64
 [Warmstart 6/10] Step 240/300 | Loss=0.1012 | PredAvg=132.6 | TargetAvg=139.8 |
 MAE=15.78
 [Warmstart 6/10] Step 260/300 | Loss=0.1023 | PredAvg=134.3 | TargetAvg=140.1 |
 MAE=16.64
 [Warmstart 6/10] Step 280/300 | Loss=0.1059 | PredAvg=133.8 | TargetAvg=140.3 |
 MAE=16.67
 [Warmstart 6/10] Step 300/300 | Loss=0.1044 | PredAvg=134.1 | TargetAvg=140.2 |
 MAE=16.55
 Warmstart epoch 6 done | Loss=0.1044
 Det eval complete | Acc=0.6500 | AvgBudget=127.8 | StdBudget=6.8 | Range=[111,
 149]
 [Warmstart 7/10] Step 20/300 | Loss=0.0147 | PredAvg=130.2 | TargetAvg=132.0 |
 MAE=7.20
 [Warmstart 7/10] Step 40/300 | Loss=0.0247 | PredAvg=131.1 | TargetAvg=134.8 |
 MAE=10.03
 [Warmstart 7/10] Step 60/300 | Loss=0.0493 | PredAvg=130.7 | TargetAvg=135.7 |

MAE=11.63
 [Warmstart 7/10] Step 80/300 | Loss=0.0476 | PredAvg=133.3 | TargetAvg=135.9 |
 MAE=12.30
 [Warmstart 7/10] Step 100/300 | Loss=0.0744 | PredAvg=133.0 | TargetAvg=138.2 |
 MAE=14.53
 [Warmstart 7/10] Step 120/300 | Loss=0.0957 | PredAvg=132.7 | TargetAvg=139.2 |
 MAE=15.43
 [Warmstart 7/10] Step 140/300 | Loss=0.0995 | PredAvg=132.5 | TargetAvg=139.2 |
 MAE=15.44
 [Warmstart 7/10] Step 160/300 | Loss=0.0883 | PredAvg=132.1 | TargetAvg=138.5 |
 MAE=14.54
 [Warmstart 7/10] Step 180/300 | Loss=0.1017 | PredAvg=132.9 | TargetAvg=140.1 |
 MAE=15.69
 [Warmstart 7/10] Step 200/300 | Loss=0.0930 | PredAvg=132.8 | TargetAvg=139.5 |
 MAE=14.88
 [Warmstart 7/10] Step 220/300 | Loss=0.0979 | PredAvg=132.5 | TargetAvg=139.9 |
 MAE=15.27
 [Warmstart 7/10] Step 240/300 | Loss=0.1017 | PredAvg=132.7 | TargetAvg=140.5 |
 MAE=15.81
 [Warmstart 7/10] Step 260/300 | Loss=0.1006 | PredAvg=132.7 | TargetAvg=140.2 |
 MAE=15.60
 [Warmstart 7/10] Step 280/300 | Loss=0.1023 | PredAvg=132.8 | TargetAvg=140.3 |
 MAE=15.68
 [Warmstart 7/10] Step 300/300 | Loss=0.1009 | PredAvg=132.6 | TargetAvg=140.2 |
 MAE=15.55
 Warmstart epoch 7 done | Loss=0.1009
 Det eval complete | Acc=0.6900 | AvgBudget=142.0 | StdBudget=6.8 | Range=[129, 161]
 Saved best warmstart by reward = 0.7915
 [Warmstart 8/10] Step 20/300 | Loss=0.0750 | PredAvg=149.0 | TargetAvg=147.8 |
 MAE=20.90
 [Warmstart 8/10] Step 40/300 | Loss=0.0582 | PredAvg=138.0 | TargetAvg=141.5 |
 MAE=16.18
 [Warmstart 8/10] Step 60/300 | Loss=0.0493 | PredAvg=136.0 | TargetAvg=138.9 |
 MAE=14.18
 [Warmstart 8/10] Step 80/300 | Loss=0.0424 | PredAvg=135.1 | TargetAvg=137.8 |
 MAE=12.99
 [Warmstart 8/10] Step 100/300 | Loss=0.0434 | PredAvg=134.8 | TargetAvg=138.0 |
 MAE=13.76
 [Warmstart 8/10] Step 120/300 | Loss=0.0810 | PredAvg=134.5 | TargetAvg=140.1 |
 MAE=15.48
 [Warmstart 8/10] Step 140/300 | Loss=0.0867 | PredAvg=134.3 | TargetAvg=140.3 |
 MAE=15.69
 [Warmstart 8/10] Step 160/300 | Loss=0.0870 | PredAvg=134.1 | TargetAvg=140.1 |
 MAE=15.44
 [Warmstart 8/10] Step 180/300 | Loss=0.0851 | PredAvg=133.6 | TargetAvg=139.7 |
 MAE=14.83
 [Warmstart 8/10] Step 200/300 | Loss=0.0804 | PredAvg=133.6 | TargetAvg=139.1 |

MAE=14.54
 [Warmstart 8/10] Step 220/300 | Loss=0.0836 | PredAvg=133.3 | TargetAvg=139.3 |
 MAE=14.84
 [Warmstart 8/10] Step 240/300 | Loss=0.0865 | PredAvg=133.4 | TargetAvg=139.4 |
 MAE=14.75
 [Warmstart 8/10] Step 260/300 | Loss=0.0848 | PredAvg=133.6 | TargetAvg=139.6 |
 MAE=14.88
 [Warmstart 8/10] Step 280/300 | Loss=0.1006 | PredAvg=133.5 | TargetAvg=140.2 |
 MAE=15.60
 [Warmstart 8/10] Step 300/300 | Loss=0.0991 | PredAvg=133.6 | TargetAvg=140.2 |
 MAE=15.44
 Warmstart epoch 8 done | Loss=0.0991
 Det eval complete | Acc=0.6500 | AvgBudget=126.1 | StdBudget=6.0 | Range=[114,
 144]
 [Warmstart 9/10] Step 20/300 | Loss=0.0864 | PredAvg=129.1 | TargetAvg=140.0 |
 MAE=14.40
 [Warmstart 9/10] Step 40/300 | Loss=0.0653 | PredAvg=129.5 | TargetAvg=138.4 |
 MAE=12.65
 [Warmstart 9/10] Step 60/300 | Loss=0.1253 | PredAvg=131.7 | TargetAvg=143.4 |
 MAE=16.53
 [Warmstart 9/10] Step 80/300 | Loss=0.1153 | PredAvg=134.0 | TargetAvg=144.1 |
 MAE=16.90
 [Warmstart 9/10] Step 100/300 | Loss=0.0989 | PredAvg=134.1 | TargetAvg=142.6 |
 MAE=16.24
 [Warmstart 9/10] Step 120/300 | Loss=0.0987 | PredAvg=133.8 | TargetAvg=141.9 |
 MAE=16.32
 [Warmstart 9/10] Step 140/300 | Loss=0.0975 | PredAvg=134.4 | TargetAvg=142.4 |
 MAE=16.56
 [Warmstart 9/10] Step 160/300 | Loss=0.1004 | PredAvg=134.1 | TargetAvg=142.5 |
 MAE=16.41
 [Warmstart 9/10] Step 180/300 | Loss=0.1101 | PredAvg=134.5 | TargetAvg=142.9 |
 MAE=16.88
 [Warmstart 9/10] Step 200/300 | Loss=0.1137 | PredAvg=134.1 | TargetAvg=142.8 |
 MAE=16.80
 [Warmstart 9/10] Step 220/300 | Loss=0.1067 | PredAvg=133.9 | TargetAvg=141.9 |
 MAE=16.19
 [Warmstart 9/10] Step 240/300 | Loss=0.1046 | PredAvg=133.7 | TargetAvg=141.4 |
 MAE=15.76
 [Warmstart 9/10] Step 260/300 | Loss=0.0989 | PredAvg=133.2 | TargetAvg=140.7 |
 MAE=15.16
 [Warmstart 9/10] Step 280/300 | Loss=0.0935 | PredAvg=133.4 | TargetAvg=140.5 |
 MAE=14.96
 [Warmstart 9/10] Step 300/300 | Loss=0.0934 | PredAvg=133.3 | TargetAvg=140.2 |
 MAE=14.68
 Warmstart epoch 9 done | Loss=0.0934
 Det eval complete | Acc=0.6100 | AvgBudget=125.9 | StdBudget=5.8 | Range=[114,
 145]
 [Warmstart 10/10] Step 20/300 | Loss=0.0054 | PredAvg=129.8 | TargetAvg=129.6 |

```

MAE=6.80
[Warmstart 10/10] Step 40/300 | Loss=0.0468 | PredAvg=129.5 | TargetAvg=133.1 |
MAE=8.93
[Warmstart 10/10] Step 60/300 | Loss=0.0704 | PredAvg=132.6 | TargetAvg=137.5 |
MAE=12.40
[Warmstart 10/10] Step 80/300 | Loss=0.0777 | PredAvg=132.1 | TargetAvg=137.4 |
MAE=12.71
[Warmstart 10/10] Step 100/300 | Loss=0.0895 | PredAvg=131.9 | TargetAvg=138.6 |
MAE=13.26
[Warmstart 10/10] Step 120/300 | Loss=0.0923 | PredAvg=133.9 | TargetAvg=139.8 |
MAE=14.71
[Warmstart 10/10] Step 140/300 | Loss=0.0883 | PredAvg=133.7 | TargetAvg=139.3 |
MAE=14.37
[Warmstart 10/10] Step 160/300 | Loss=0.0797 | PredAvg=133.1 | TargetAvg=138.6 |
MAE=13.66
[Warmstart 10/10] Step 180/300 | Loss=0.0802 | PredAvg=133.6 | TargetAvg=138.6 |
MAE=13.86
[Warmstart 10/10] Step 200/300 | Loss=0.0787 | PredAvg=133.5 | TargetAvg=139.1 |
MAE=14.01
[Warmstart 10/10] Step 220/300 | Loss=0.0735 | PredAvg=133.4 | TargetAvg=138.6 |
MAE=13.62
[Warmstart 10/10] Step 240/300 | Loss=0.0832 | PredAvg=133.6 | TargetAvg=139.3 |
MAE=14.18
[Warmstart 10/10] Step 260/300 | Loss=0.0887 | PredAvg=133.9 | TargetAvg=140.1 |
MAE=14.98
[Warmstart 10/10] Step 280/300 | Loss=0.0915 | PredAvg=134.1 | TargetAvg=140.4 |
MAE=15.32
[Warmstart 10/10] Step 300/300 | Loss=0.0904 | PredAvg=133.9 | TargetAvg=140.2 |
MAE=15.23
Warmstart epoch 10 done | Loss=0.0904
  Det eval complete | Acc=0.6200 | AvgBudget=128.6 | StdBudget=6.5 | Range=[111,
149]

```

Loading best warm-start policy by accuracy...

Post warm-start evaluation:

```

  Det eval complete | Acc=0.6900 | AvgBudget=145.9 | StdBudget=10.7 |
Range=[117, 171]

```

```

{
  "warmstart_eval": {
    "accuracy": 0.69,
    "avg_thinking_tokens": 145.86,
    "avg_reward": 0.7895700000000001,
    "avg_budget": 145.86,
    "min_budget": 117.0,
    "max_budget": 171.0,
    "std_budget": 10.689265643625852
  }
}

```

```
}
```

Pre-computing all training features...

Pre-training value network on oracle rewards...

Target oracle avg reward: 1.1133

Value pretrain epoch 1/5 | Loss=2.1873

Value pretrain epoch 2/5 | Loss=0.3601

Value pretrain epoch 3/5 | Loss=0.2858

Value pretrain epoch 4/5 | Loss=0.2439

Value pretrain epoch 5/5 | Loss=0.3169

Value pre-training complete. PPO will start with better advantage estimates.

Starting PPO training...

PPO epsilon (clip range): 0.2

PPO epochs per rollout: 2

Batch size: 8 questions at once on GPU

Pre-encoding all training features...

```
=====
```

PPO Epoch 1/10

```
=====
```

Phase 1: Collecting rollout experiences (batched)...

Rollout done | Acc=0.7433 | AvgReward=0.8459 | AvgBudget=166.6 |
StdBudget=45.0

Phase 2: PPO updates (2 epochs, batch=32)...

PPO update epoch 1/2 | PolicyLoss=-0.0179 | ValueLoss=0.7911 |
ClipFrac=0.106

PPO update epoch 2/2 | PolicyLoss=-0.0179 | ValueLoss=0.3705 |
ClipFrac=0.000

Evaluating deterministic policy on test set...

Det eval complete | Acc=0.7300 | AvgBudget=152.1 | StdBudget=9.6 | Range=[125,
175]

Test Accuracy : 0.7300

Test AvgTokens : 152.1

Test AvgReward : 0.8365

Test Budget : 152.1 +/- 9.6 [125, 175]

Saved best PPO model by reward = 0.8365

Saved best PPO model by accuracy = 0.7300

```
=====
```

PPO Epoch 2/10

```
=====
```

Phase 1: Collecting rollout experiences (batched)...

Rollout done | Acc=0.7467 | AvgReward=0.8477 | AvgBudget=171.3 |
StdBudget=44.5

```

Phase 2: PPO updates (2 epochs, batch=32)...
  PPO update epoch 1/2 | PolicyLoss=-0.0244 | ValueLoss=0.4566 |
ClipFrac=0.131
  PPO update epoch 2/2 | PolicyLoss=0.0247 | ValueLoss=0.4129 | ClipFrac=0.081

Evaluating deterministic policy on test set...
Det eval complete | Acc=0.7400 | AvgBudget=158.3 | StdBudget=9.6 | Range=[130,
179]
Test Accuracy   : 0.7400
Test AvgTokens  : 158.3
Test AvgReward  : 0.8458
Test Budget     : 158.3 +/- 9.6 [130, 179]
Saved best PPO model by reward = 0.8458
Saved best PPO model by accuracy = 0.7400

```

```

=====
PPO Epoch 3/10
=====

```

```

Phase 1: Collecting rollout experiences (batched)...
Rollout done | Acc=0.7400 | AvgReward=0.8411 | AvgBudget=167.8 |
StdBudget=43.9
Phase 2: PPO updates (2 epochs, batch=32)...
  PPO update epoch 1/2 | PolicyLoss=-0.0241 | ValueLoss=0.6712 |
ClipFrac=0.096
  PPO update epoch 2/2 | PolicyLoss=0.0102 | ValueLoss=0.5237 | ClipFrac=0.153

Evaluating deterministic policy on test set...
Det eval complete | Acc=0.7600 | AvgBudget=162.5 | StdBudget=8.6 | Range=[137,
181]
Test Accuracy   : 0.7600
Test AvgTokens  : 162.5
Test AvgReward  : 0.8687
Test Budget     : 162.5 +/- 8.6 [137, 181]
Saved best PPO model by reward = 0.8687
Saved best PPO model by accuracy = 0.7600

```

```

=====
PPO Epoch 4/10
=====

```

```

Phase 1: Collecting rollout experiences (batched)...
Rollout done | Acc=0.7567 | AvgReward=0.8619 | AvgBudget=167.9 |
StdBudget=44.6
Phase 2: PPO updates (2 epochs, batch=32)...
  PPO update epoch 1/2 | PolicyLoss=-0.0234 | ValueLoss=0.3914 |
ClipFrac=0.021
  PPO update epoch 2/2 | PolicyLoss=0.0107 | ValueLoss=0.3840 | ClipFrac=0.159

Evaluating deterministic policy on test set...

```

Det eval complete | Acc=0.7500 | AvgBudget=170.4 | StdBudget=5.8 | Range=[154, 185]

Test Accuracy : 0.7500
Test AvgTokens : 170.4
Test AvgReward : 0.8523
Test Budget : 170.4 +/- 5.8 [154, 185]

=====
PP0 Epoch 5/10
=====

Phase 1: Collecting rollout experiences (batched)...

Rollout done | Acc=0.7400 | AvgReward=0.8393 | AvgBudget=171.5 | StdBudget=44.4

Phase 2: PP0 updates (2 epochs, batch=32)...

PP0 update epoch 1/2 | PolicyLoss=-0.0260 | ValueLoss=0.3788 | ClipFrac=0.000

PP0 update epoch 2/2 | PolicyLoss=-0.0243 | ValueLoss=0.3505 | ClipFrac=0.061

Evaluating deterministic policy on test set...

Det eval complete | Acc=0.7800 | AvgBudget=176.9 | StdBudget=3.7 | Range=[167, 187]

Test Accuracy : 0.7800
Test AvgTokens : 176.9
Test AvgReward : 0.8866
Test Budget : 176.9 +/- 3.7 [167, 187]
Saved best PP0 model by reward = 0.8866
Saved best PP0 model by accuracy = 0.7800

=====
PP0 Epoch 6/10
=====

Phase 1: Collecting rollout experiences (batched)...

Rollout done | Acc=0.7500 | AvgReward=0.8509 | AvgBudget=173.2 | StdBudget=47.1

Phase 2: PP0 updates (2 epochs, batch=32)...

PP0 update epoch 1/2 | PolicyLoss=-0.0114 | ValueLoss=0.3413 | ClipFrac=0.035

PP0 update epoch 2/2 | PolicyLoss=-0.0156 | ValueLoss=0.3059 | ClipFrac=0.091

Evaluating deterministic policy on test set...

Det eval complete | Acc=0.7600 | AvgBudget=182.2 | StdBudget=3.4 | Range=[172, 191]

Test Accuracy : 0.7600
Test AvgTokens : 182.2
Test AvgReward : 0.8589
Test Budget : 182.2 +/- 3.4 [172, 191]

=====
PPO Epoch 7/10
=====

Phase 1: Collecting rollout experiences (batched)...
Rollout done | Acc=0.7667 | AvgReward=0.8704 | AvgBudget=175.9 |
StdBudget=46.0
Phase 2: PPO updates (2 epochs, batch=32)...
PPO update epoch 1/2 | PolicyLoss=-0.0362 | ValueLoss=0.3606 |
ClipFrac=0.008
PPO update epoch 2/2 | PolicyLoss=0.0093 | ValueLoss=0.2990 | ClipFrac=0.087

Evaluating deterministic policy on test set...
Det eval complete | Acc=0.8000 | AvgBudget=184.7 | StdBudget=3.1 | Range=[175,
192]
Test Accuracy : 0.8000
Test AvgTokens : 184.7
Test AvgReward : 0.9077
Test Budget : 184.7 +/- 3.1 [175, 192]
Saved best PPO model by reward = 0.9077
Saved best PPO model by accuracy = 0.8000

=====
PPO Epoch 8/10
=====

Phase 1: Collecting rollout experiences (batched)...
Rollout done | Acc=0.7733 | AvgReward=0.8768 | AvgBudget=179.8 |
StdBudget=47.7
Phase 2: PPO updates (2 epochs, batch=32)...
PPO update epoch 1/2 | PolicyLoss=0.0147 | ValueLoss=0.3248 | ClipFrac=0.015
PPO update epoch 2/2 | PolicyLoss=-0.0285 | ValueLoss=0.2999 |
ClipFrac=0.034

Evaluating deterministic policy on test set...
Det eval complete | Acc=0.7600 | AvgBudget=186.6 | StdBudget=2.8 | Range=[180,
193]
Test Accuracy : 0.7600
Test AvgTokens : 186.6
Test AvgReward : 0.8567
Test Budget : 186.6 +/- 2.8 [180, 193]

=====
PPO Epoch 9/10
=====

Phase 1: Collecting rollout experiences (batched)...
Rollout done | Acc=0.7733 | AvgReward=0.8753 | AvgBudget=182.8 |
StdBudget=44.0
Phase 2: PPO updates (2 epochs, batch=32)...

PP0 update epoch 1/2 | PolicyLoss=-0.0112 | ValueLoss=0.2662 |
ClipFrac=0.019

PP0 update epoch 2/2 | PolicyLoss=-0.0187 | ValueLoss=0.2772 |
ClipFrac=0.136

Evaluating deterministic policy on test set...

Det eval complete | Acc=0.7700 | AvgBudget=193.1 | StdBudget=2.9 | Range=[186,
201]

Test Accuracy : 0.7700

Test AvgTokens : 193.1

Test AvgReward : 0.8659

Test Budget : 193.1 +/- 2.9 [186, 201]

=====

PP0 Epoch 10/10

=====

Phase 1: Collecting rollout experiences (batched)...

Rollout done | Acc=0.7833 | AvgReward=0.8855 | AvgBudget=187.3 |
StdBudget=45.1

Phase 2: PP0 updates (2 epochs, batch=32)...

PP0 update epoch 1/2 | PolicyLoss=0.0040 | ValueLoss=0.2529 | ClipFrac=0.000

PP0 update epoch 2/2 | PolicyLoss=0.0055 | ValueLoss=0.2684 | ClipFrac=0.028

Evaluating deterministic policy on test set...

Det eval complete | Acc=0.7800 | AvgBudget=194.0 | StdBudget=6.2 | Range=[181,
210]

Test Accuracy : 0.7800

Test AvgTokens : 194.0

Test AvgReward : 0.8780

Test Budget : 194.0 +/- 6.2 [181, 210]

Loading best PP0 model for final evaluation...

Final deterministic policy evaluation:

Det eval complete | Acc=0.8000 | AvgBudget=184.7 | StdBudget=3.1 | Range=[175,
192]

```
{
  "final_deterministic_eval": {
    "accuracy": 0.8,
    "avg_thinking_tokens": 184.68,
    "avg_reward": 0.9076599999999999,
    "avg_budget": 184.68,
    "min_budget": 175.0,
    "max_budget": 192.0,
    "std_budget": 3.136494858914964
  }
}
```

Saved: final_policy_v15.pt
Saved: final_value_v15.pt

=====			
FINAL RESULTS COMPARISON			
=====			
Method	Acc	AvgTokens	AvgReward

Fixed budget 96	0.4700	96	0.5395
Fixed budget 128	0.6300	128	0.7235
Fixed budget 160	0.7400	160	0.8450
Fixed budget 192	0.7600	192	0.8540
Fixed budget 224	0.7900	224	0.8755
Fixed budget 256	0.8400	256	0.9220
Policy (PPO deterministic)	0.8000	185	0.9077
=====			

AI Disclosure: Written with Claude's assistance

[]: